

IBM HR Analytics Employee Attrition & Performance

```
In [1]: # importing libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# to display plots and graphs inline
%matplotlib inline
```

```
In [2]: # using the seaborn style for graphs
plt.style.use("seaborn-v0_8")
```

```
In [3]: # read the dataset
employee_data = pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
In [4]: employee_data.shape
```

```
Out[4]: (1470, 35)
```

```
In [5]: employee_data.head()
```

```
Out[5]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educa
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

```
In [6]: # finding missing values
employee_data.isnull().sum()
```

```
Out[6]: Age 0
Attrition 0
BusinessTravel 0
DailyRate 0
Department 0
DistanceFromHome 0
Education 0
EducationField 0
EmployeeCount 0
EmployeeNumber 0
EnvironmentSatisfaction 0
Gender 0
HourlyRate 0
JobInvolvement 0
JobLevel 0
JobRole 0
JobSatisfaction 0
MaritalStatus 0
MonthlyIncome 0
MonthlyRate 0
NumCompaniesWorked 0
Over18 0
OverTime 0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
StandardHours 0
StockOptionLevel 0
TotalWorkingYears 0
TrainingTimesLastYear 0
WorkLifeBalance 0
YearsAtCompany 0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

```
In [7]: #checking data types
employee_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                      1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

Exploratory Data Analysis


```

In [8]: # basic descriptive statistics
employee_data.describe().T

```

Out[8]:

	count	mean	std	min	25%	50%	75%	
Age	1470.0	36.923810	9.135373	18.0	30.00	36.0	43.00	
DailyRate	1470.0	802.485714	403.509100	102.0	465.00	802.0	1157.00	1
DistanceFromHome	1470.0	9.192517	8.106864	1.0	2.00	7.0	14.00	
Education	1470.0	2.912925	1.024165	1.0	2.00	3.0	4.00	
EmployeeCount	1470.0	1.000000	0.000000	1.0	1.00	1.0	1.00	
EmployeeNumber	1470.0	1024.865306	602.024335	1.0	491.25	1020.5	1555.75	2
EnvironmentSatisfaction	1470.0	2.721769	1.093082	1.0	2.00	3.0	4.00	
HourlyRate	1470.0	65.891156	20.329428	30.0	48.00	66.0	83.75	
JobInvolvement	1470.0	2.729932	0.711561	1.0	2.00	3.0	3.00	
JobLevel	1470.0	2.063946	1.106940	1.0	1.00	2.0	3.00	
JobSatisfaction	1470.0	2.728571	1.102846	1.0	2.00	3.0	4.00	
MonthlyIncome	1470.0	6502.931293	4707.956783	1009.0	2911.00	4919.0	8379.00	19
MonthlyRate	1470.0	14313.103401	7117.786044	2094.0	8047.00	14235.5	20461.50	26
NumCompaniesWorked	1470.0	2.693197	2.498009	0.0	1.00	2.0	4.00	
PercentSalaryHike	1470.0	15.209524	3.659938	11.0	12.00	14.0	18.00	
PerformanceRating	1470.0	3.153741	0.360824	3.0	3.00	3.0	3.00	
RelationshipSatisfaction	1470.0	2.712245	1.081209	1.0	2.00	3.0	4.00	
StandardHours	1470.0	80.000000	0.000000	80.0	80.00	80.0	80.00	
StockOptionLevel	1470.0	0.793878	0.852077	0.0	0.00	1.0	1.00	
TotalWorkingYears	1470.0	11.279592	7.780782	0.0	6.00	10.0	15.00	
TrainingTimesLastYear	1470.0	2.799320	1.289271	0.0	2.00	3.0	3.00	
WorkLifeBalance	1470.0	2.761224	0.706476	1.0	2.00	3.0	3.00	
YearsAtCompany	1470.0	7.008163	6.126525	0.0	3.00	5.0	9.00	
YearsInCurrentRole	1470.0	4.229252	3.623137	0.0	2.00	3.0	7.00	
YearsSinceLastPromotion	1470.0	2.187755	3.222430	0.0	0.00	1.0	3.00	
YearsWithCurrManager	1470.0	4.123129	3.568136	0.0	2.00	3.0	7.00	



In [9]:

```
# mapping the attrition 1 => yes and 0 => no in the new column
employee_data["left"] = np.where(employee_data["Attrition"] == "Yes",1,0)
```

In [10]:

```
employee_data.head()
```

```
Out[10]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educa
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 36 columns

```
In [11]: # supressing all the warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [12]: def Numerical_Variables_TargetPlots(df,segment_by,target_var ="Attrition"):
        """A function for plotting the distribution of numerical variables
        and its effect on attrition """

        fig, ax = plt.subplots(ncols= 2, figsize = (14,6))

        #boxplot for comparison
        sns.boxplot(x= target_var, y= segment_by, data=df,ax=ax[0])
        ax[0].set_title("Comparison of " + segment_by + " vs " + target_var)

        #distribution plot
        ax[1].set_title("Distribution of "+segment_by)
        ax[1].set_ylabel("Frequency")
        sns.distplot(a= df[segment_by],ax=ax[1], kde=False)

        plt.show()
```

```
In [13]: def Categorical_Variables_targetPlots(df, segment_by,invert_axis =False,target_var
        """A function for plotting the effect of variables(categorical data)
        on Attrition"""

        fig, ax = plt.subplots(ncols=2, figsize = (14,6))

        # countplot for distribution along with target variable
        # invert axis variable helps to interchange the axis so that
        # names of categories doesn't overlap
        if invert_axis == False:
            sns.countplot(x = segment_by, data=df, hue="Attrition",ax=ax[0])
        else:
            sns.countplot(y = segment_by, data=df, hue="Attrition",ax=ax[0])

        ax[0].set_title("Comparison of " + segment_by + " vs " + "Attrition")

        #plot the effect of the variable on attrition
        if invert_axis == False:
            sns.barplot(x = segment_by, y = target_var, data=df,ci=None)
        else:
            sns.barplot(y= segment_by, x = target_var, data=df, ci=None)
```

```
ax[1].set_title("Attrition rate by {}".format(segment_by))
ax[1].set_ylabel("Average(Attrition)")
plt.tight_layout()


plt.show()
```

Analyzing the variables

- Numerical Variables

Age



In [14]: *# we are checking the distribution of employee age and its related to attrition or*
 Numerical_Variables_TargetPlots(employee_data, segment_by="Age")





- **Age Distribution:** Employees who have stayed with the company ("No" attrition status) tend to be older on average compared to those who have left ("Yes" attrition status).
- **Median Age:** The median age for employees who have left is 30, while for those who stayed it is 37. This suggests that younger employees are more likely to leave the company.
- **Quartiles:** The age range between the first quartile and the third quartile is wider for employees who stayed, indicating a more diverse age range in the retained employees.
- **Outliers:** The presence of outliers in the "Yes" category might suggest that while the majority of younger employees are more likely to leave, there are a few older employees who also leave the company.

Daily Rate & Monthly Income & HourlyRate

In [15]: *#Analyzing the daily wage rate vs employee left the company or not*
 Numerical_Variables_TargetPlots(employee_data, "DailyRate")





```
In [16]: Numerical_Variables_TargetPlots(employee_data, "MonthlyIncome")
```



- Employee's working with lower daily rates are more prone to leave the company than compared to the employee's working with higher rates. The same trend is resonated with monthly income too.

Hourly Rate



```
In [17]: Numerical_Variables_TargetPlots(employee_data, "HourlyRate")
```



- From plot we have seen that there is no significant difference in the hourly rate and attrition. Therefore hourly rate is considered as not significant to attrition.

Percent Salary Hike


```
In [18]: Numerical_Variables_TargetPlots(employee_data, "PercentSalaryHike")
```



- Majority (60% of total strength) of employee's receive 16% salary hike in the company, employee's who received less salary hike have left the company.

Total Working Years

```
In [19]: Numerical_Variables_TargetPlots(employee_data, "TotalWorkingYears")
```



```
In [20]: sns.lmplot(x="TotalWorkingYears", y="PercentSalaryHike", data=employee_data, fit_
             height=6, aspect=1.5)
plt.show()
```



- Employee's with less working years have received 25% Salary hike when they switch to another company, but there is no linear relationship between working years and salary hike.
- Attrition is not seen among the employee's having more than 20 years of experience if their salary hike is more than 20%, even if the salary hike is below 20% attrition rate among the employee's is very low.
- Employee's with lesser years of experience are prone to leave the company in search of better pay, irrespective of salary hike

Distance From Home

```
In [21]: Numerical_Variables_TargetPlots(employee_data, "DistanceFromHome")
```



- There is a higher number of people who reside near to offices and hence the attrition levels are lower for distances less than 10. With increase in distance from home, attrition rate also increases

Analysing the Variables

- Categorical Variables

Job Involvement

```
In [22]: #cross tabulation between attrition and Job Involvement
pd.crosstab(employee_data.JobInvolvement, employee_data.Attrition)
```

Out[22]:


	Attrition	No	Yes
JobInvolvement			
1	55	28	
2	304	71	
3	743	125	
4	131	13	

```
In [23]: # calculate the percentage of people having different job involvement role
round(employee_data.JobInvolvement.value_counts()/employee_data.shape[0]*100,2)
```

Out[23]:

```
JobInvolvement
3    59.05
2    25.51
4     9.80
1     5.65
Name: count, dtype: float64
```


```
In [24]: Categorical_Variables_targetPlots(employee_data, "JobInvolvement")
```



1. In the total data set, 59% have high job involvement whereas 25% have medium involvement rate
2. From above plot we can observe that round 50% of people in low job involvement (level 1 & 2) have left the company.
3. Even the people who have high job involvement have higher attrition rate around 15% in that category have left company

Job Satisfaction

In [25]: `Categorical_Variables_targetPlots(employee_data, "JobSatisfaction")`



As expected, people with low satisfaction have left the company around 23% in that category. what surprising is out of the people who rated medium and high job satisfaction around 32% has left the company. There should be some other factor which triggers their exit from the company

Performance Rating

In [26]: `# Checking the number of categories under performance rating
employee_data.PerformanceRating.value_counts()`


```
Out[26]: PerformanceRating
3      1244
4       226
Name: count, dtype: int64
```

```
In [27]: # calculate the percentage of performance rating per category in the whole dataset
round(employee_data.PerformanceRating.value_counts()/employee_data.shape[0]*100,2)
```

```
Out[27]: PerformanceRating
3      84.63
4      15.37
Name: count, dtype: float64
```

- Around 85% of the people in the company rated as Excellent and remaining 15% rated as Outstanding

```
In [28]: Categorical_Variables_targetPlots(employee_data,"PerformanceRating")
```



Contrary to normal belief that employee's having higher rating will not leave the company. It may be seen that there is no significant difference between the Performance rating and Attrition rate


WorkLifeBalance

```
In [29]: # percentage of worklife balance rating accross the company date
round(employee_data.WorkLifeBalance.value_counts()/employee_data.shape[0]*100,2)
```

```
Out[29]: WorkLifeBalance
3      60.75
2      23.40
4      10.41
1       5.44
Name: count, dtype: float64
```


More than 60% of the employee's rated that they have Better worklife balance and 10% rated for Best worklife balance

```
In [30]: Categorical_Variables_targetPlots(employee_data,"WorkLifeBalance")
```



- As expected more than 30% of the people who rated as Bad WorkLifeBalance have left the company and around 15% of the people who rated for Best WorkLifeBalance also left the company


In [31]: `Categorical_Variables_targetPlots(employee_data,"OverTime")`



- More than 30% of employee's who worked overtime has left the company, where as 90% of employee's who have not experienced overtime has not left the company. Therefore overtime is a strong indicator of attrition

Business Travel

In [32]: `Categorical_Variables_targetPlots(employee_data,segment_by="BusinessTravel")`



- There are more people who travel rarely compared to people who travel frequently. In case of people who travel Frequency around 25% of people have left the company and in other cases attrition rate doesn't vary significantly on travel

Department


In [33]: `employee_data.Department.value_counts()`

Out[33]:

Department	count
Research & Development	961
Sales	446
Human Resources	63

Name: count, dtype: int64

In [34]: `Categorical_Variables_targetPlots(employee_data,segment_by="Department")`




- On comparing departmentwise, we can conclude that HR has seen only a marginal high in turnover rates whereas the numbers are significant in sales department with turnover rates of 39 %. The attrition levels are not appreciable in R & D where 67 % have recorded no attrition.
- Sales has seen higher attrition levels about 20.6% followed by HR around 18%

Education Field


```
In [35]: employee_data.EducationField.value_counts()
```

```
Out[35]: EducationField
Life Sciences      606
Medical            464
Marketing           159
Technical Degree    132
Other               82
Human Resources     27
Name: count, dtype: int64
```

```
In [36]: Categorical_Variables_targetPlots(employee_data,"EducationField",invert_axis=True)
```



```
In [37]: plt.figure(figsize=(10,8))
sns.barplot(y = "EducationField", x = "left", hue = "Education", data = employee_data)
plt.show()
```




- There are more people with a Life sciences followed by medical and marketing
- Employee's in the EducationField of Human Resources and Technical Degree have highest attrition levels around 26% and 23% respectively

- When compared with Education level, we have observed that employees in the highest level of education in there field of study have left the company. We can conclude that EducationField is a strong indicator of attrition

Environment Satisfaction


In [38]: `Categorical_Variables_targetPlots(employee_data,"EnvironmentSatisfaction")`




we can see that people having low environment satisfaction 25% leave the company

Gender Vs Attrition

In [39]: `sns.boxplot(x=employee_data['Gender'], y=employee_data["MonthlyIncome"])
plt.title("Monthly Income Vs Gender Boxplot", fontsize=20)
plt.xlabel("Gender", fontsize=16)
plt.ylabel('Monthly Income', fontsize=16)
plt.show()`




```
In [40]: Categorical_Variables_targetPlots(employee_data, "Gender")
```



- Monthly Income distribution for Male and Female is almost similar, so the attrition rate of Male and Female is almost around 15%. Gender is not a strong indicator of attrition

```
In [41]: fig, ax = plt.subplots(2,3, figsize=(20,20)) # "ax" has references to all the four
plt.suptitle("Comparison of various factors vs Gender", fontsize = 20)
sns.barplot(x = employee_data["Gender"], y = employee_data["DistanceFromHome"], hue=
sns.barplot(x = employee_data["Gender"], y = employee_data["YearsAtCompany"], hue=
sns.barplot(x = employee_data["Gender"], y = employee_data["TotalWorkingYears"], hu
sns.barplot(x = employee_data["Gender"], y = employee_data["YearsInCurrentRole"], h
sns.barplot(x = employee_data["Gender"], y = employee_data["YearsSinceLastPromotion"], h
sns.barplot(x = employee_data["Gender"], y = employee_data["NumCompaniesWorked"], h
plt.show()
```


Comparison of various factors vs Gender



1. Distance from home matters to women employees more than men.
2. Female employees are spending more years in one company compare to their counterpart.
3. Female employees spending more years in current company are more inclined to switch.

Job Role


```
In [42]: Categorical_Variables_targetPlots(employee_data, "JobRole", invert_axis=True)
```



1. Jobs held by the employee is maximum in Sales Executive, then R&D , then Laboratory Technician
2. People working in Sales department is most likely quit the company followed by Laboratory Technician and Human Resources there attrition rates are 40%, 24% and 22% respectively

Marital Status

In [43]: `Categorical_Variables_targetPlots(employee_data,"MaritalStatus")`



From the plot,it is understood that irrespective of the marital status,there are large people who stay with the company and do not leave.Therefore,marital status is a weak predictor of attrition

In []:

Building Decision Tree

```
In [44]: from sklearn.model_selection import train_test_split


# for fitting classification tree
from sklearn.tree import DecisionTreeClassifier

# to create a confusion matrix
```

```
from sklearn.metrics import confusion_matrix

# import whole class of metrics
from sklearn import metrics
```

```
In [45]: employee_data.Attrition.value_counts().plot(kind="bar")
plt.xlabel("Attrition")
plt.ylabel("Count")
plt.show()
```



```
In [46]: employee_data["Attrition"].value_counts()
```

```
Out[46]: Attrition
No      1233
Yes      237
Name: count, dtype: int64
```

From the Exploratory data analysis, variable that are not significant to attrition are:

- EmployeeCount, EmployeeNumber, Gender, HourlyRate, JobLevel, MaritalStatus, Over18, StandardHours

```
In [47]: # copying the main employee data to another dataframe
employee_data_new = employee_data.copy()
```

```
In [48]: # dropping the not significant variables
employee_data_new.drop(["EmployeeCount", "EmployeeNumber", "Gender", "HourlyRate", "Over18", "StandardHours", "JobLevel", "MaritalStatus"], axis=1, inplace=True)
```

```
In [49]: employee_data_new.head()
```

Out[49]:	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educa
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 29 columns

Handling Categorical Variables

- Segregate the numerical and Categorical variables
- Convert Categorical variables to dummy variables

```
In [50]: # data types of variables
dict(employee_data_new.dtypes)
```

```
Out[50]: {'Age': dtype('int64'),
'Attrition': dtype('O'),
'BusinessTravel': dtype('O'),
'DailyRate': dtype('int64'),
'Department': dtype('O'),
'DistanceFromHome': dtype('int64'),
'Education': dtype('int64'),
'EducationField': dtype('O'),
'EnvironmentSatisfaction': dtype('int64'),
'JobInvolvement': dtype('int64'),
'JobLevel': dtype('int64'),
'JobRole': dtype('O'),
'JobSatisfaction': dtype('int64'),
'MaritalStatus': dtype('O'),
'MonthlyIncome': dtype('int64'),
'MonthlyRate': dtype('int64'),
'NumCompaniesWorked': dtype('int64'),
'OverTime': dtype('O'),
'PercentSalaryHike': dtype('int64'),
'PerformanceRating': dtype('int64'),
'RelationshipSatisfaction': dtype('int64'),
'StockOptionLevel': dtype('int64'),
'TotalWorkingYears': dtype('int64'),
'TrainingTimesLastYear': dtype('int64'),
'WorkLifeBalance': dtype('int64'),
'YearsAtCompany': dtype('int64'),
'YearsInCurrentRole': dtype('int64'),
'YearsSinceLastPromotion': dtype('int64'),
'YearsWithCurrManager': dtype('int64')}
```

```
In [51]: #segregating the variables based on datatypes

numeric_variable_names = [key for key in dict(employee_data_new.dtypes) if dict(en
categorical_variable_names = [key for key in dict(employee_data_new.dtypes) if dict
```

```
In [52]: numeric_variable_names
```

```
Out[52]: ['Age',
'DailyRate',
'DistanceFromHome',
'Education',
'EnvironmentSatisfaction',
'JobInvolvement',
'JobLevel',
'JobSatisfaction',
'MonthlyIncome',
'MonthlyRate',
'NumCompaniesWorked',
'PercentSalaryHike',
'PerformanceRating',
'RelationshipSatisfaction',
'StockOptionLevel',
'TotalWorkingYears',
'TrainingTimesLastYear',
'WorkLifeBalance',
'YearsAtCompany',
'YearsInCurrentRole',
'YearsSinceLastPromotion',
'YearsWithCurrManager']
```

```
In [53]: categorical_variable_names
```

```
Out[53]: ['Attrition',
'BusinessTravel',
'Department',
'EducationField',
'JobRole',
'MaritalStatus',
'Overtime']
```

```
In [54]: # store the numerical variables data in separate dataset

employee_data_num = employee_data_new[numeric_variable_names]
```

```
In [55]: # store the categorical variables data in separate dataset

employee_data_cat = employee_data_new[categorical_variable_names]

#dropping the attrition
employee_data_cat.drop(["Attrition"], axis=1, inplace=True)
```

```
In [56]: # converting into dummy variables

employee_data_cat = pd.get_dummies(employee_data_cat)
```

```
In [57]: # merging the both numerical and categorical data


employee_data_final = pd.concat([employee_data_num, employee_data_cat, employee_data
```

```
In [58]: employee_data_final.head()
```

Out[58]:

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	JobInvolvement	JobL
0	41	1102		1	2	2	3
1	49	279		8	1	3	2
2	37	1373		2	2	4	2
3	33	1392		3	4	4	3
4	27	591		2	1	1	3

5 rows × 49 columns



```
In [59]: #final features
features = list(employee_data_final.columns.difference(["Attrition"]))
```

```
In [60]: features
```

```
Out[60]: ['Age',
          'BusinessTravel_Non-Travel',
          'BusinessTravel_Travel_Frequently',
          'BusinessTravel_Travel_Rarely',
          'DailyRate',
          'Department_Human Resources',
          'Department_Research & Development',
          'Department_Sales',
          'DistanceFromHome',
          'Education',
          'EducationField_Human Resources',
          'EducationField_Life Sciences',
          'EducationField_Marketing',
          'EducationField_Medical',
          'EducationField_Other',
          'EducationField_Technical Degree',
          'EnvironmentSatisfaction',
          'JobInvolvement',
          'JobLevel',
          'JobRole_Healthcare Representative',
          'JobRole_Human Resources',
          'JobRole_Laboratory Technician',
          'JobRole_Manager',
          'JobRole_Manufacturing Director',
          'JobRole_Research Director',
          'JobRole_Research Scientist',
          'JobRole_Sales Executive',
          'JobRole_Sales Representative',
          'JobSatisfaction',
          'MaritalStatus_Divorced',
          'MaritalStatus_Married',
          'MaritalStatus_Single',
          'MonthlyIncome',
          'MonthlyRate',
          'NumCompaniesWorked',
          'OverTime_No',
          'OverTime_Yes',
          'PercentSalaryHike',
          'PerformanceRating',
          'RelationshipSatisfaction',
          'StockOptionLevel',
          'TotalWorkingYears',
          'TrainingTimesLastYear',
          'WorkLifeBalance',
          'YearsAtCompany',
          'YearsInCurrentRole',
          'YearsSinceLastPromotion',
          'YearsWithCurrManager']
```

```
In [ ]:
```

Separating the Target and the Predictors

```
In [61]: # separating the target and predictors
```

```
X = employee_data_final[features]
y = employee_data_final[["Attrition"]]
```

```
In [62]: X.shape
```

```
Out[62]: (1470, 48)
```


Train_Test Split (Stratified Sampling of Y)

```
In [63]: # Function for creating model pipelines
from sklearn.pipeline import make_pipeline

# function for crossvalidate score
from sklearn.model_selection import cross_validate

# to find the best
from sklearn.model_selection import GridSearchCV
```


```
In [64]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size= 0.3, stratify= y)
```

```
In [65]: # Checks
#proportion in training_data
y_train.Attrition.value_counts()/len(y_train)
```

```
Out[65]: Attrition
No      0.838678
Yes     0.161322
Name: count, dtype: float64
```

```
In [66]: #Checks
#Proportion in training data
pd.DataFrame(y_train.Attrition.value_counts()/len(y_train)).plot(kind = "bar")
```

```
Out[66]: <Axes: xlabel='Attrition'>
```



```
In [67]: #Proportion of test data
y_test.Attrition.value_counts()/len(y_test)
```

```
Out[67]: Attrition
No      0.839002
Yes     0.160998
Name: count, dtype: float64
```

```
In [68]: #make a pipeline for the decision tree model

pipelines = {
    "clf" : make_pipeline(DecisionTreeClassifier(max_depth=3, random_state=100))
}
```

Cross Validate

- To check the accuracy of the pipeline

```
In [69]: scores = cross_validate(pipelines['clf'], X_train, y_train, return_train_score=True)
```

```
In [70]: scores['test_score'].mean()
```

```
Out[70]: np.float64(0.8396590101823348)
```

Average accuracy of pipeline with Decision Tree Classifier is 83.48%

Cross-Validation and Hyper Parameters Tuning

Cross Validation is the process of finding the best combination of parameters for the model by training and evaluating the model for each combination of the parameters

- Declare a hyper-parameters to fine tune the Decision Tree Classifier

Decision Tree is a greedy algorithm it searches the entire space of possible decision trees. so we need to find a optimum parameter(s) or criteria for stopping the decision tree at some point. We use the hyperparameters to prune the decision tree

```
In [71]: decisiontree_hyperparameters = {
    "decisiontreeclassifier__max_depth": np.arange(3,12),
    "decisiontreeclassifier__max_features": np.arange(3,10),
    "decisiontreeclassifier__min_samples_split": [2,3,4,5,6,7,8,9,10,11,12,13,14,15],
    "decisiontreeclassifier__min_samples_leaf" : np.arange(1,3)
}
```

```
In [72]: pipelines['clf']
```

```
Out[72]: Pipeline
└── DecisionTreeClassifier
```

Decision Tree Classifier with gini index

Fit and tune models with cross-validation


Now that we have our `pipelines` and `hyperparameters` dictionaries declared, we're ready to tune our models with cross-validation.

- We are doing 5 fold cross validation

```
In [73]: #Create a cross validation object from decision tree classifier and it's hyperparameters
clf_model = GridSearchCV(pipelines['clf'], decisiontree_hyperparameters, cv=5, n_jobs=5)
```

```
In [74]: #fit the model with train data
clf_model.fit(X_train, y_train)
```

```
Out[74]:
```



```
In [75]: #Display the best parameters for the Decision Tree Model
clf_model.best_params_
```

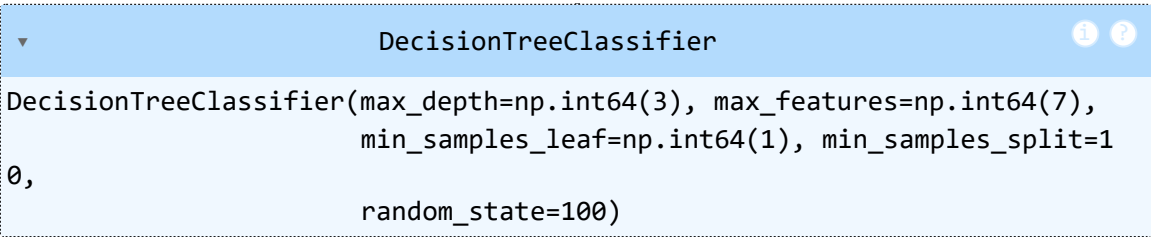
```
Out[75]: {'decisiontreeclassifier__max_depth': np.int64(3),
          'decisiontreeclassifier__max_features': np.int64(7),
          'decisiontreeclassifier__min_samples_leaf': np.int64(1),
          'decisiontreeclassifier__min_samples_split': 10}
```

```
In [76]: #Display the best score for the fitted model
clf_model.best_score_
```

```
Out[76]: np.float64(0.8561780724603363)
```

```
In [77]: #In Pipeline we can use the string names to get the decisiontree classifier
clf_model.best_estimator_.named_steps["decisiontreeclassifier"]
```

```
Out[77]:
```



```
In [78]: # saving into a variable to get graph
clf_best_model = clf_model.best_estimator_.named_steps["decisiontreeclassifier"]
```

Model Performance Evaluation

- On Test Data

```
In [79]: #Making a dataframe of actual and predicted data from test set

tree_test_pred = pd.concat([y_test.reset_index(drop=True),pd.DataFrame(clf_model.pr
tree_test_pred.columns = ["actual","predicted"]

#setting the index to original index
tree_test_pred.index = y_test.index
```

```
In [80]: tree_test_pred.head()
```

```
Out[80]:
```

	actual	predicted
34	Yes	Yes
1432	No	No
334	No	No
1068	Yes	No
736	No	No

```
In [81]: #keeping only positive condition (yes for attrition)

pred_probability = pd.DataFrame(p[1] for p in clf_model.predict_proba(X_test))
pred_probability.columns = ["predicted_prob"]
pred_probability.index = y_test.index
```

```
In [82]: # merging the predicted data and it's probability value

tree_test_pred = pd.concat([tree_test_pred,pred_probability],axis=1)
```

```
In [83]: tree_test_pred.head()
```

```
Out[83]:
```

	actual	predicted	predicted_prob
34	Yes	Yes	0.632184
1432	No	No	0.220859
334	No	No	0.072165
1068	Yes	No	0.145985
736	No	No	0.072165

```
In [84]: #converting the Labels Yes --> 1 and No --> 0 for further operations below

tree_test_pred["actual_left"] = np.where(tree_test_pred["actual"] == "Yes",1,0)
tree_test_pred["predicted_left"] = np.where(tree_test_pred["predicted"] == "Yes",1,
```

```
In [85]: tree_test_pred.head()
```

```
Out[85]:
```

	actual	predicted	predicted_prob	actual_left	predicted_left
34	Yes	Yes	0.632184	1	1
1432	No	No	0.220859	0	0
334	No	No	0.072165	0	0
1068	Yes	No	0.145985	1	0
736	No	No	0.072165	0	0

Confusion Matrix

The confusion matrix is a way of tabulating the number of misclassifications, i.e., the number of predicted classes which ended up in a wrong classification bin based on the true classes.


```
In [86]: #confusion matrix
metrics.confusion_matrix(tree_test_pred.actual, tree_test_pred.predicted, labels=["Yes", "No"])
```

```
Out[86]: array([[ 21,  50],
               [ 26, 344]])
```

```
In [87]: #confusion matrix visualisation using seaborn heatmap

sns.heatmap(metrics.confusion_matrix(tree_test_pred.actual, tree_test_pred.predicted,
                                     labels=["Yes", "No"]), cmap="Greens", annot=True,
            xticklabels= ["Left", "Not Left"], yticklabels= ["Left", "Not Left"])

plt.ylabel("True label")
plt.xlabel("Predicted label")
plt.show()
```



```
In [88]: # Area Under ROC Curve
auc_score_test = metrics.roc_auc_score(tree_test_pred.actual_left, tree_test_pred.pr
```


```
print("AUROC Score :",round(auc_score_test,4))
```

AUROC Score : 0.6128

In [89]: *##Plotting the ROC Curve*

```
fpr, tpr, thresholds = metrics.roc_curve(tree_test_pred,actual_left, tree_test_pred)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label='ROC curve (area = %0.4f)' % auc_score_test)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic cuve')
plt.legend(loc="lower right")
plt.show()
```



From the ROC Curve, we have a choice to make depending on the value we place on true positive and tolerance for false positive rate

- If we wish to find the more people who are leaving, we could increase the true positive rate by adjusting the probability cutoff for classification. However by doing so would also increase the false positive rate. we need to find the optimum value of cutoff for classification

Metrics

- Recall: Ratio of the total number of correctly classified positive examples divide to the total number of positive examples. High Recall indicates the class is correctly recognized

- Precision: To get the value of precision we divide the total number of correctly classified positive examples by the total number of predicted positive examples. High Precision indicates an example labeled as positive is indeed positive

```
In [90]: #calculating the recall score

print("Recall Score:",round(metrics.recall_score(tree_test_pred.actual_left,tree_te

Recall Score: 29.577
```

```
In [91]: #calculating the precision score

print("Precision Score:",round(metrics.precision_score(tree_test_pred.actual_left,t

Precision Score: 44.681
```

```
In [92]: print(metrics.classification_report(tree_test_pred.actual_left,tree_test_pred.predi

          precision    recall  f1-score   support

         0         0.87        0.93        0.90        370
         1         0.45        0.30        0.36         71

 accuracy                   0.83        441
 macro avg              0.66        0.61        0.63        441
 weighted avg           0.80        0.83        0.81        441
```

Visualization of Decision Tree

- Dependencies
 - Need to install graphviz (conda install pydot graphviz)
 - Set the environment path variable to graphviz folder

```
In [93]: # conda install pydot graphviz
         # pip install pydotplus
```

```
In [94]: from sklearn.tree import export_graphviz
```

```
In [95]: import pydotplus as plot
```

```
In [96]: from six import StringIO
         from IPython.display import Image
         from sklearn.tree import export_graphviz
         import pydotplus as pdot
```

```
In [97]: import os
         os.environ["PATH"] += os.pathsep + "C:/Program Files/Graphviz/bin/"
```

```
In [98]: #write the dot data
         dot_data = StringIO()
```

```
In [99]: #export the decision tree along with the feature names into a dot file format


         export_graphviz(clf_best_model,out_file=dot_data,filled=True,
                        rounded=True,special_characters=True,feature_names = X_train.columr
```

```
In [100... #make a graph from dot file

         graph = pdot.graph_from_dot_data(dot_data.getvalue())
```

```
In [101... Image(graph.create_png())
```

```
Out[101]:
```



```
In [102... #export the tree diagram
graph.write_png("Employee_Attrition.png")
```

```
Out[102]: True
```

Conclusions

- The dependent variable for the decision tree is a OverTime which has two classes Yes or No.
- The most influential attribute to determine how to classify employee will leave or not is the Monthly Income variable.
- Even though the employee's having the average monthly income if their job involvement is average or less than average, those employees will most likely leave the company.
- Employee's whose monthly income is less than 2000 approximately will leave the company even if their total working years is more than 6 months.

Suggested Actions:

It's not sensible to focus on every employee who wants to leave because it costs time and energy for human resources management department. HR department need to focus on:

- Improving the work conditions

Provide an option for the employee's to work from home, on a flexible schedule, or in an office with an ergonomic workspace, they will be more satisfied with their work and more likely to achieve a healthy work-life balance.

- Offer modest salaries and perks

To maintain the critical employee's company need's to offer equitable and modest salaries. You can also give added perks like flexible schedules, travel discounts etc.

- Employee Engagement

When you have talented employee's we need to find ways that you can help expand the employee's skill set, so that their involvement in the job increases. If their involvement is low, they will get bored and think that they are not growing within the organization

In []:

In []: